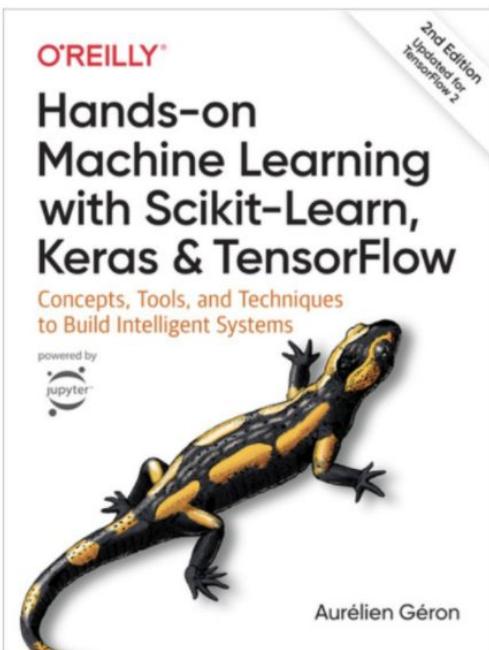


Deep Learning Adventures + San Diego Machine Learning

Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow, 2nd Edition

★★★★★ [92 reviews](#)

By [Aurélien Géron](#)



TIME TO COMPLETE:

24h 18m

TOPICS:

[Machine Learning](#)

PUBLISHED BY:

[O'Reilly Media, Inc.](#)

PUBLICATION DATE:

September 2019

PRINT LENGTH:

848 pages

Chapter Title

1	The Machine Learning Landscape
2	End-to-End Machine Learning Project
3	Classification
4	Training Models
5	Support Vector Machines
6	Decision Trees
7	Ensemble Learning and Random Forests
8	Dimensionality Reduction
9a	Unsupervised Learning Techniques: Clustering
9b	Unsupervised Learning Techniques: Gaussian Mixtures
10	Introduction to Artificial Neural Networks with Keras
11	Training Deep Neural Networks
12	Custom Models and Training with TensorFlow
13	Loading and Preprocessing Data with TensorFlow
14	Deep Computer Vision Using Convolutional Neural Networks
15	Processing Sequences Using RNNs and CNNs
16	Natural Language Processing with RNNs and Attention
17	Representation Learning and Generative Learning Using Autoencoders and GANs
18	Reinforcement Learning
19	Training and Deploying TensorFlow Models at Scale

Chapter 14 : Deep Computer Vision Using Convolutional Neural Networks
Discussion led by George Zoto



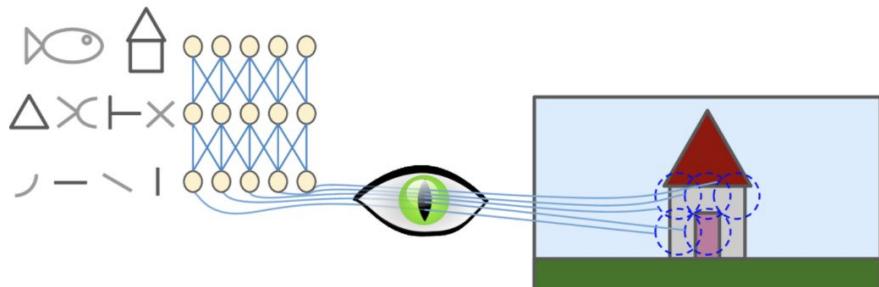


Figure 14-1. Biological neurons in the visual cortex respond to specific patterns in small regions of the visual field called receptive fields; as the visual signal makes its way through consecutive brain modules, neurons respond to more complex patterns in larger receptive fields.

Vertical edge detection

3 ⁽¹⁾	0 ⁽⁰⁾	1 ⁽¹⁾	2	7	4	
1 ⁽¹⁾	5 ⁽⁰⁾	8 ⁽⁴⁾	9	3	1	
→2 ⁽¹⁾	7 ⁽⁰⁾	2 ⁽¹⁾	5	1	3	
0	1	3	1	7	8	
4	2	1	6	2	8	
2	4	5	2	3	9	

6x6

"convolution"

*

1	0	-1	
1	0	-1	
1	0	-1	

3x3

filter

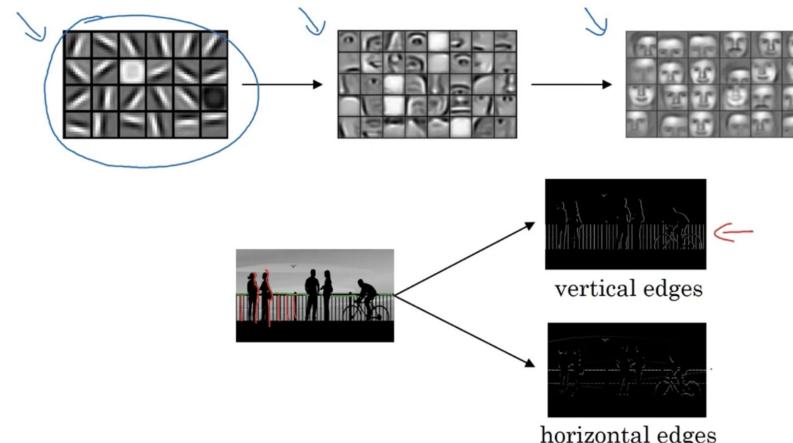
=

-5			

4x4

Andrew Ng

Computer Vision Problem



Andrew Ng

Vertical edge detection

$$\begin{array}{|c|c|c|c|c|c|} \hline
 10 & 10 & 10 & 0 & 0 & 0 \\ \hline
 10 & 10 & 10 & 0 & 0 & 0 \\ \hline
 10 & 10 & 10 & 0 & 0 & 0 \\ \hline
 10 & \underline{10} & \underline{10} & \underline{0} & 0 & 0 \\ \hline
 10 & \underline{10} & \underline{10} & \underline{0} & 0 & 0 \\ \hline
 10 & \underline{10} & \underline{10} & \underline{0} & 0 & 0 \\ \hline
 \end{array} \quad * \quad
 \begin{array}{|c|c|c|} \hline
 1 & 0 & -1 \\ \hline
 1 & 0 & -1 \\ \hline
 1 & 0 & -1 \\ \hline
 \end{array} \quad = \quad
 \begin{array}{|c|c|c|c|} \hline
 0 & 30 & 30 & 0 \\ \hline
 0 & 30 & 30 & 0 \\ \hline
 0 & 30 & 30 & 0 \\ \hline
 0 & \boxed{30} & 30 & 0 \\ \hline
 \end{array}$$

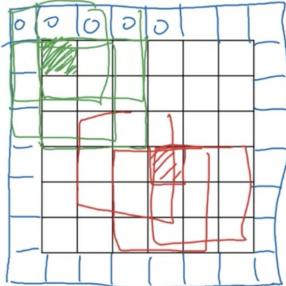
3×3

↑ 4x4

Andrew

Andrew Ng

Padding



$$\frac{6 \times 6}{n \times n} \rightarrow 8 \times 8$$

$$P = \text{padding} = 1$$

- shrinks output
- throws away info from edge

$$3 \times 3 \\ f \times f$$

=

$$4 \times 4$$

$$n-f+1 \times n-f+1$$

$$6-3+1=4$$

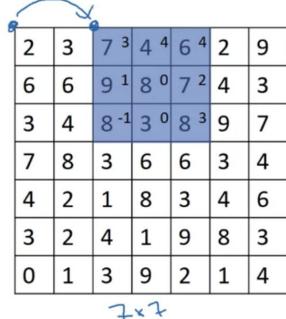
$$n+2p-f+1 \times n+2p-f+1$$

$$6+2-3+1 \times$$

$$= 6 \times 6$$

Andrew Ng

Strided convolution



*

$$3 \quad 4 \quad 4 \\ 1 \quad 0 \quad 2 \\ -1 \quad 0 \quad 3$$

=

$$3 \times 3$$

$$\text{stride} = 2$$

Valid and Same convolutions

→ no padding

$$\begin{array}{lll} \text{"Valid":} & n \times n & \times \quad f \times f \\ & 6 \times 6 & \times \quad 3 \times 3 \end{array} \rightarrow \begin{array}{l} \underline{n-f+1} \times \underline{n-f+1} \\ 4 \times 4 \end{array}$$

"Same": Pad so that output size is the same as the input size.

$$\begin{array}{l} n+2p-f+1 \times n+2p-f+1 \\ n+2p-f+1 = n \Rightarrow p = \frac{f-1}{2} \end{array}$$

Andrew Ng

Summary of convolutions

$$n \times n \text{ image} \quad f \times f \text{ filter}$$

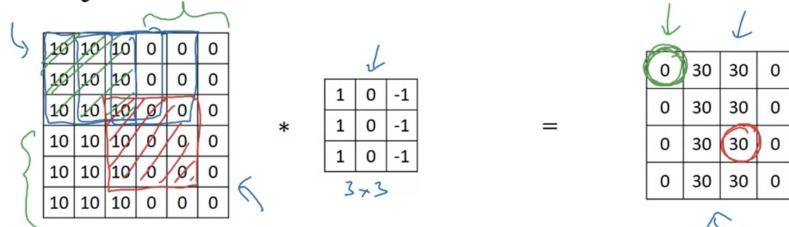
$$\text{padding } p \quad \text{stride } s$$

$$\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \quad \times \quad \left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor$$

Andrew Ng

Enhancing Vision with Convolutional Neural Networks

Why convolutions

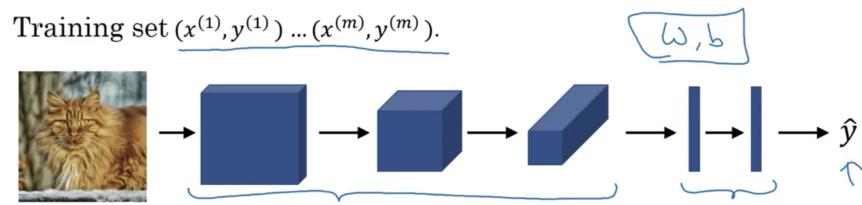


Parameter sharing: A feature detector (such as a vertical edge detector) that's useful in one part of the image is probably useful in another part of the image.

→ **Sparsity of connections:** In each layer, each output value depends only on a small number of inputs.

Putting it together

Training set $(x^{(1)}, y^{(1)}) \dots (x^{(m)}, y^{(m)})$.

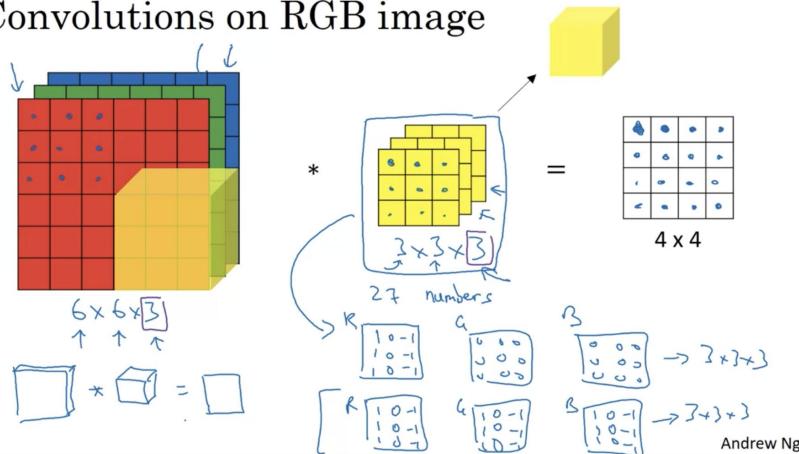


$$\text{Cost } J = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

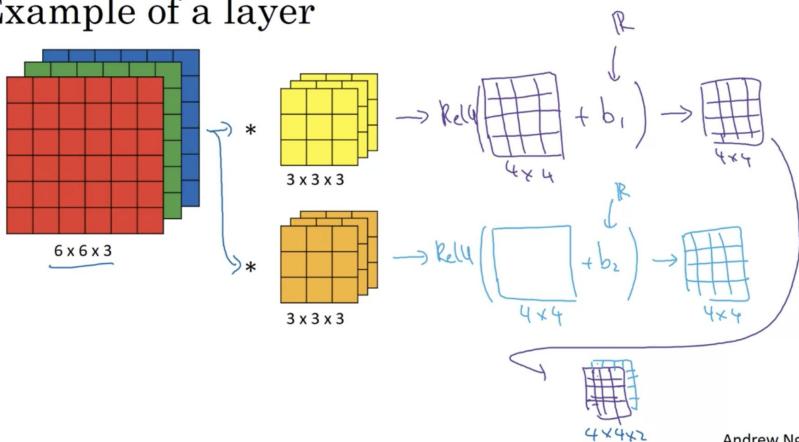
Use gradient descent to optimize parameters to reduce J

Content added after Week 1 and 2, based on group feedback

Convolutions on RGB image



Example of a layer



Input Volume (+pad 1) ($7 \times 7 \times 3$)	Filter W0 ($3 \times 3 \times 3$)	Filter W1 ($3 \times 3 \times 3$)	Output Volume ($3 \times 3 \times 2$)
$x[:, :, 0]$	$w0[:, :, 0]$	$w1[:, :, 0]$	$o[:, :, 0]$
$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 2 & 0 \\ 0 & 1 & 0 & 2 & 0 & 1 & 0 \\ 0 & 1 & 0 & 2 & 2 & 0 & 0 \\ 0 & 2 & 0 & 0 & 2 & 0 & 0 \\ 0 & 2 & 1 & 2 & 2 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$	$\begin{bmatrix} -1 & 0 & 1 \\ 0 & 0 & 1 \\ 1 & -1 & 1 \end{bmatrix}$	$\begin{bmatrix} 0 & 1 & -1 \\ 0 & -1 & 0 \\ 0 & -1 & 1 \end{bmatrix}$	$\begin{bmatrix} 2 & 3 & 3 \\ 3 & 7 & 3 \\ 8 & 10 & -3 \end{bmatrix}$
$x[:, :, 1]$	$w0[:, :, 1]$	$w1[:, :, 1]$	$o[:, :, 1]$
$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 2 & 1 & 2 & 1 & 1 & 0 \\ 0 & 2 & 1 & 2 & 0 & 1 & 0 \\ 0 & 0 & 2 & 1 & 0 & 1 & 0 \\ 0 & 1 & 2 & 2 & 2 & 2 & 0 \\ 0 & 0 & 1 & 2 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$	$\begin{bmatrix} -1 & 0 & 1 \\ 1 & -1 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	$\begin{bmatrix} -1 & 0 & 0 \\ 1 & -1 & 0 \\ 1 & -1 & 0 \end{bmatrix}$	$\begin{bmatrix} -8 & -8 & -3 \\ -3 & 1 & 0 \\ -3 & -8 & -5 \end{bmatrix}$
$x[:, :, 2]$	$w0[:, :, 2]$	$w1[:, :, 2]$	$o[:, :, 2]$
$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 2 & 1 & 2 & 0 & 1 & 0 \\ 0 & 2 & 1 & 2 & 0 & 1 & 0 \\ 0 & 0 & 2 & 1 & 0 & 1 & 0 \\ 0 & 1 & 2 & 2 & 2 & 2 & 0 \\ 0 & 0 & 1 & 2 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$	$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 0 \\ 0 & -1 & 0 \end{bmatrix}$	$\begin{bmatrix} -1 & 1 & -1 \\ 0 & -1 & -1 \\ 1 & 0 & 0 \end{bmatrix}$	$\begin{bmatrix} 1 & 0 & 0 \\ 0 & -1 & -1 \\ 1 & 0 & 0 \end{bmatrix}$
Bias $b0 (1 \times 1 \times 1)$	$b0[:, :, 0]$		
0	1		

$x[:, :, 2]$
$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 2 & 1 & 1 & 2 & 0 & 0 \\ 0 & 1 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 2 & 1 & 0 & 0 \\ 0 & 2 & 2 & 1 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$
$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 2 & 1 & 1 & 2 & 0 & 0 \\ 0 & 1 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 2 & 1 & 0 & 0 \\ 0 & 2 & 2 & 1 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$
$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 2 & 1 & 1 & 2 & 0 & 0 \\ 0 & 1 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 2 & 1 & 0 & 0 \\ 0 & 2 & 2 & 1 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$
$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 2 & 1 & 1 & 2 & 0 & 0 \\ 0 & 1 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 2 & 1 & 0 & 0 \\ 0 & 2 & 2 & 1 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$

Source: <https://www.coursera.org/learn/convolutional-neural-networks>
<https://images.app.goo.gl/Z4xDvji3Gcerpm27>

toggle movement

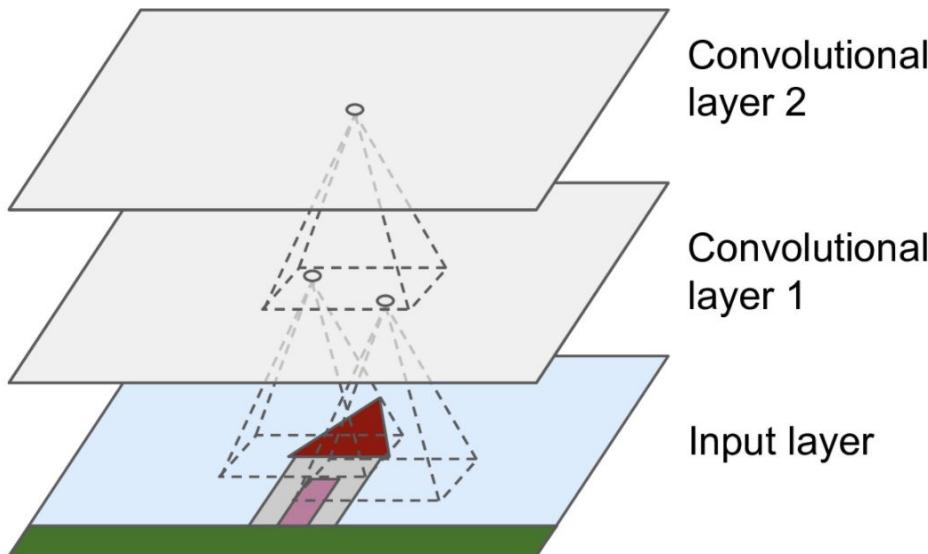


Figure 14-2. CNN layers with rectangular local receptive fields

```
model = keras.models.Sequential([
    keras.layers.Conv2D(64, 7, activation="relu", padding="same",
                       input_shape=[28, 28, 1]),
    keras.layers.MaxPooling2D(2),
    keras.layers.Conv2D(128, 3, activation="relu", padding="same"),
    keras.layers.Conv2D(128, 3, activation="relu", padding="same"),
    keras.layers.MaxPooling2D(2),
    keras.layers.Conv2D(256, 3, activation="relu", padding="same"),
    keras.layers.Conv2D(256, 3, activation="relu", padding="same"),
    keras.layers.MaxPooling2D(2),
    keras.layers.Flatten(),
    keras.layers.Dense(128, activation="relu"),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(64, activation="relu"),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(10, activation="softmax")
])
```

```
model.layers
```

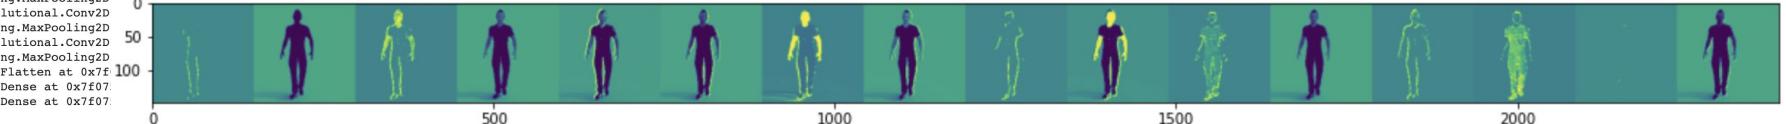
- valhuman01-00.png(image/png) -:
Saving valhuman01-00.png to
[1.]

[1.]

```
successive_outputs = [layer.output for layer in model.layers[1:]]
```

```
model.layers[0].Conv2D
```

max_pooling2d



conv2d_1



max_pooling2d_1



conv2d_2



max_pooling2d_2



conv2d_3



max_pooling2d_3



conv2d_4



max_pooling2d_4



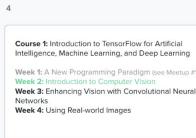
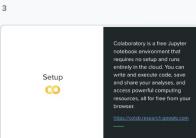
Layer (type)	Output Shape
conv2d (Conv2D)	(None, 298, 298, 16)
max_pooling2d (MaxPooling2D)	(None, 149, 149, 16)
conv2d_1 (Conv2D)	(None, 147, 147, 32)
	73, 73, 32)
	71, 71, 64)
	35, 35, 64)
	33, 33, 64)
	16, 16, 64)
	14, 14, 64)
	7, 7, 64)
	3136)
	512)
	1)

Deep Learning Adventures TensorFlow In Practice - Presentation 2

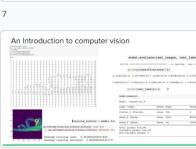
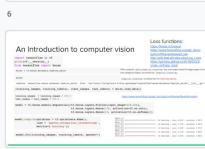
A quick overview of our course's TensorFlow In Practice specialization area.
Robert Krueg, David Petrov, George Zoto
<http://bit.ly/dla-youtube>



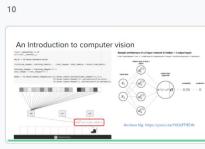
1



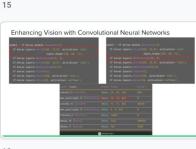
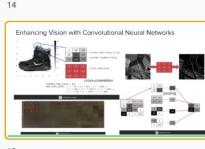
5



9



13



14

15

16

Deep-Learning-Adventures-Tensorflow-In-Practice-Presentation-2

<http://bit.ly/dla-youtube>

Introduction to Deep Learning



TensorFlow



PLAY ALL

Deep Learning Adventures - Tensorflow In Practice Specialization

13 videos • 6,564 views • Last updated on Jul 10, 2020

Public ▾



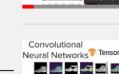
Join our Deep Learning Adventures
community and become an expert in Deep
Learning, TensorFlow, Computer Vision,
Convolutional Neural Networks, Kaggle

SORT



Introduction to Deep Learning and TensorFlow

George Zoto



Introduction to Computer Vision and Convolutional Neural Networks

George Zoto



Convolutional Neural Networks and ImageDataGenerator

George Zoto



Kaggle Challenge, Data Augmentation and Dropouts

George Zoto



Transfer Learning, Multiclass Classifications and Overfitting

George Zoto

Deep-Learning-Adventures-Tensorflow-In-Practice-Presentation-3

<http://bit.ly/dla-youtube>

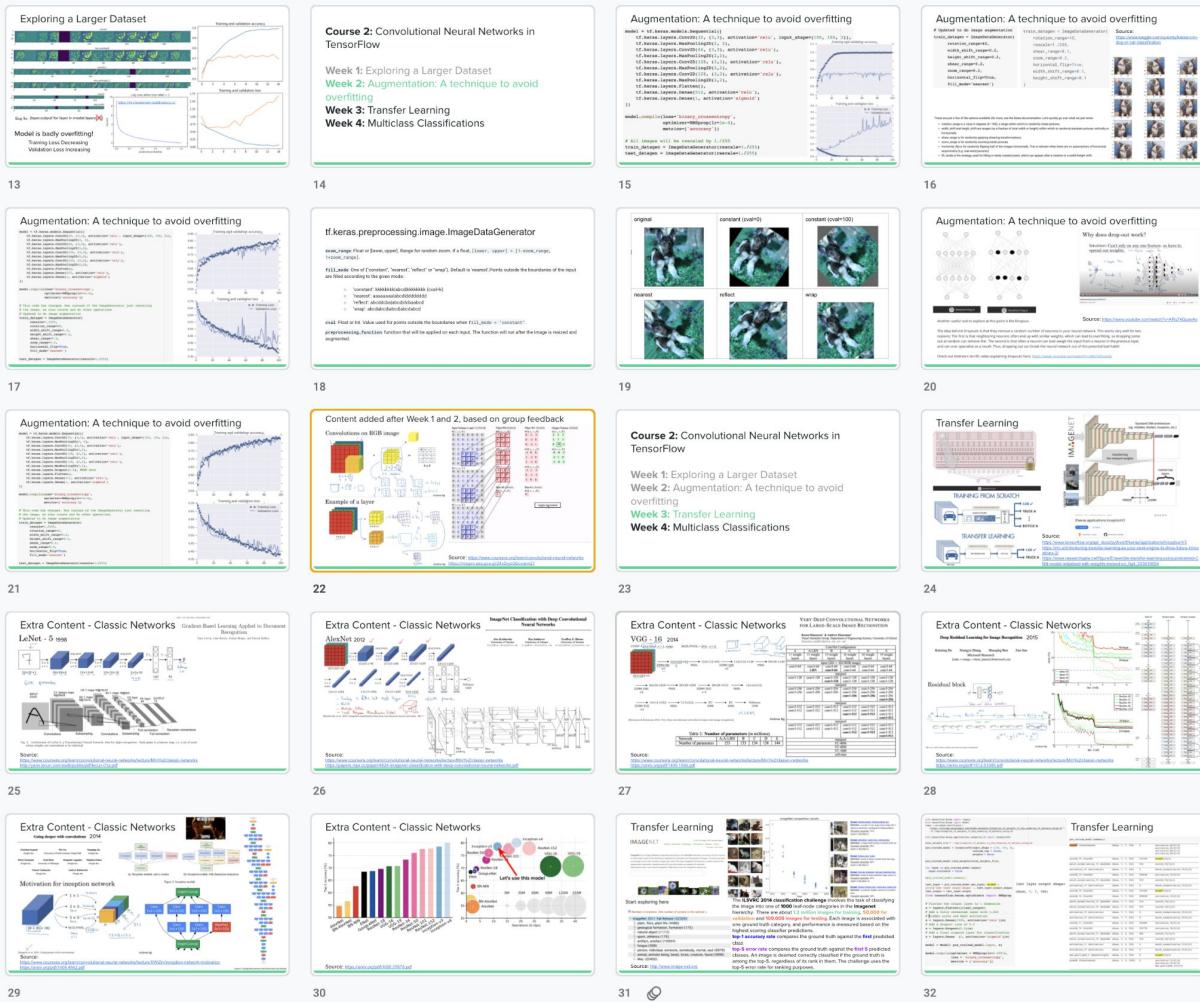


Table 14-1. LeNet-5 architecture

Layer	Type	Maps	Size	Kernel size	Stride	Activation
Out	Fully connected	–	10	–	–	RBF
F6	Fully connected	–	84	–	–	tanh
C5	Convolution	120	1×1	5×5	1	tanh
S4	Avg pooling	16	5×5	2×2	2	tanh
C3	Convolution	16	10×10	5×5	1	tanh
S2	Avg pooling	6	14×14	2×2	2	tanh
C1	Convolution	6	28×28	5×5	1	tanh
In	Input	1	32×32	–	–	–

Table 14-2. AlexNet architecture

Layer	Type	Maps	Size	Kernel size	Stride	Padding	Activation
Out	Fully connected	–	1,000	–	–	–	Softmax
F10	Fully connected	–	4,096	–	–	–	ReLU
F9	Fully connected	–	4,096	–	–	–	ReLU
S8	Max pooling	256	6×6	3×3	2	valid	–
C7	Convolution	256	13×13	3×3	1	same	ReLU
C6	Convolution	384	13×13	3×3	1	same	ReLU
C5	Convolution	384	13×13	3×3	1	same	ReLU
S4	Max pooling	256	13×13	3×3	2	valid	–
C3	Convolution	256	27×27	5×5	1	same	ReLU
S2	Max pooling	96	27×27	3×3	2	valid	–
C1	Convolution	96	55×55	11×11	4	valid	ReLU
In	Input	3 (RGB)	227×227	–	–	–	–

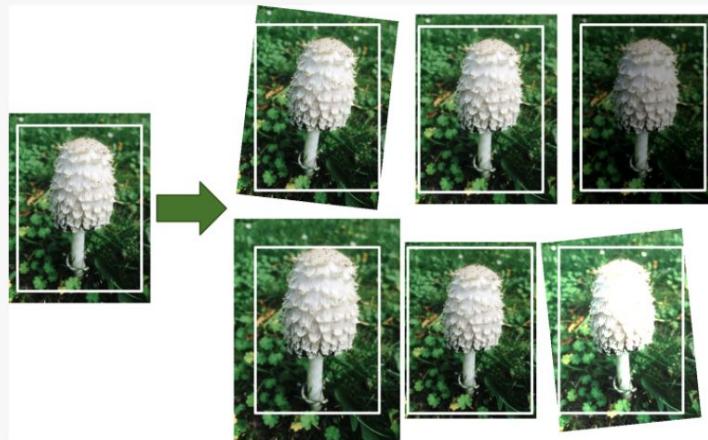


Figure 14-12. Generating new training instances from existing ones

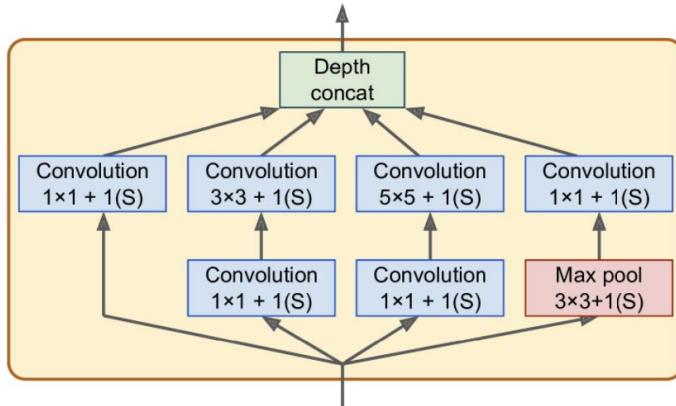


Figure 14-13. Inception module

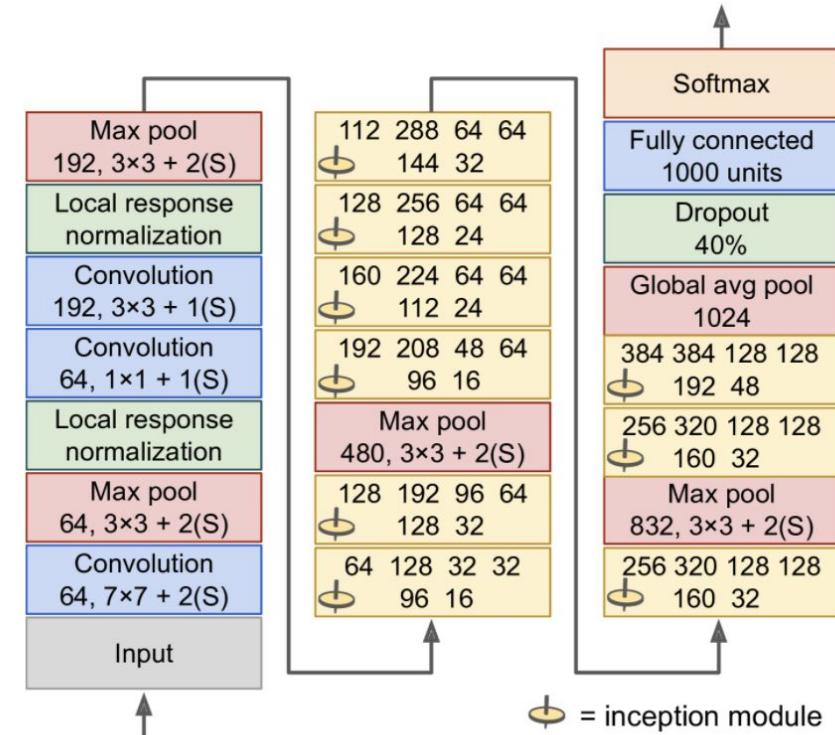


Figure 14-14. GoogLeNet architecture

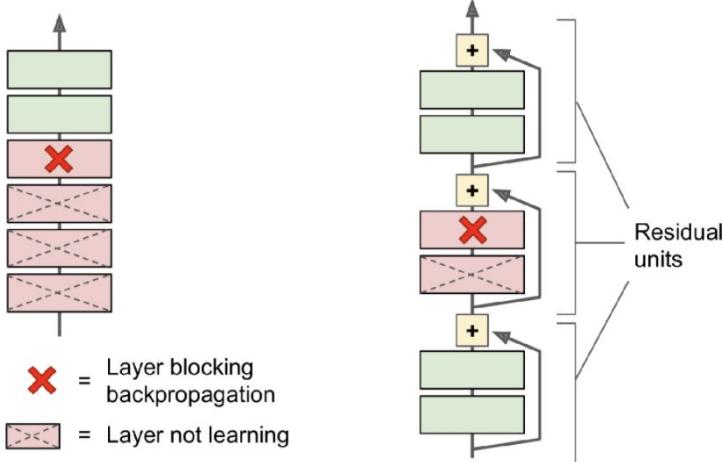


Figure 14-16. Regular deep neural network (left) and deep residual network (right)

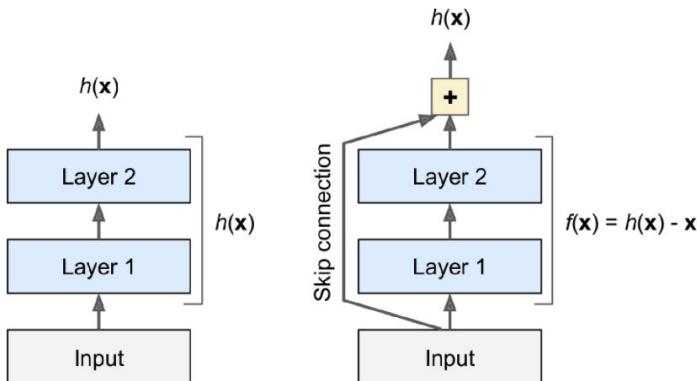


Figure 14-15. Residual learning

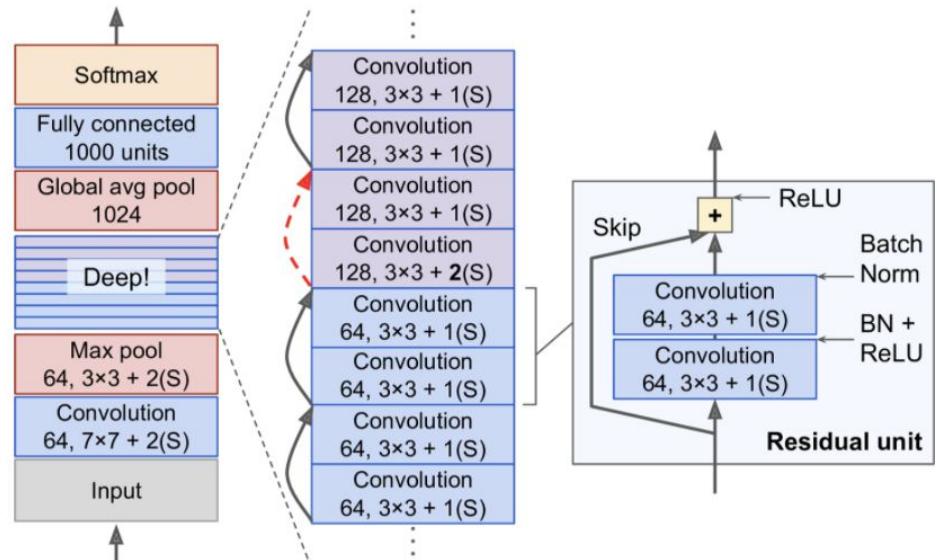


Figure 14-17. ResNet architecture

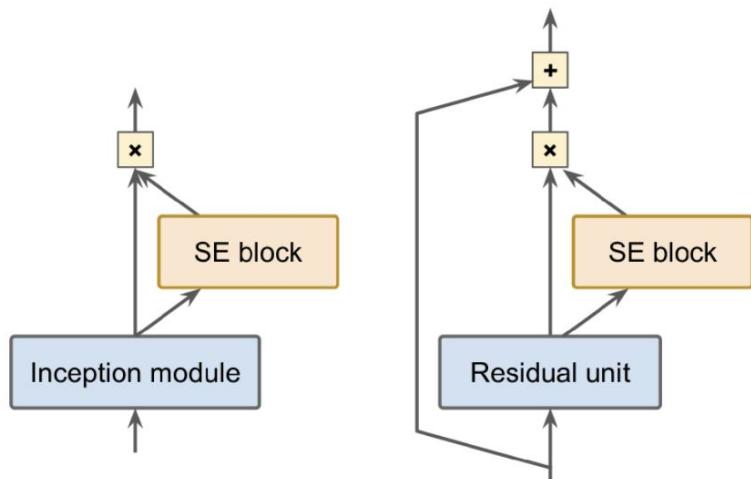


Figure 14-20. SE-Inception module (left) and SE-ResNet unit (right)

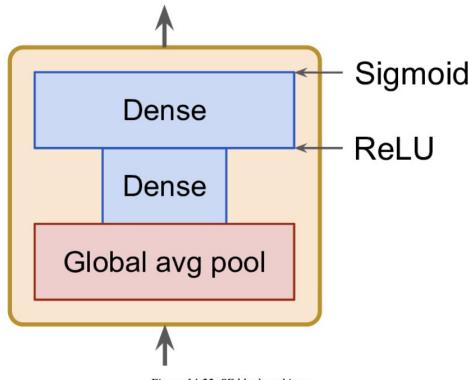


Figure 14-22. SE block architecture

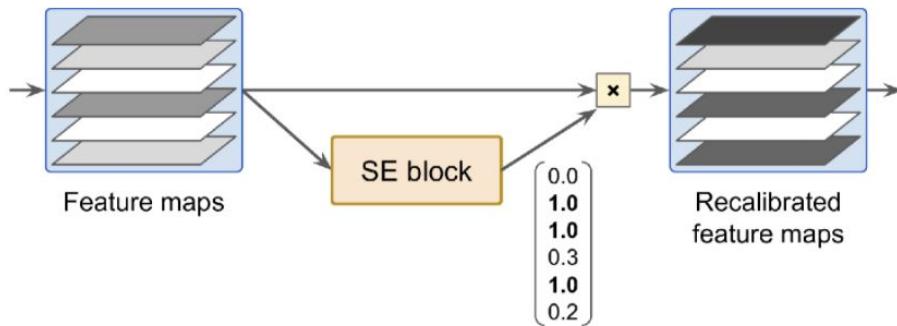


Figure 14-21. An SE block performs feature map recalibration

```

model = keras.applications.resnet50.ResNet50(weights="imagenet")

base_model = keras.applications.xception.Xception(weights="imagenet",
                                                   include_top=False)
avg = keras.layers.GlobalAveragePooling2D()(base_model.output)
output = keras.layers.Dense(n_classes, activation="softmax")(avg)
model = keras.Model(inputs=base_model.input, outputs=output)

optimizer = keras.optimizers.SGD(lr=0.2, momentum=0.9, decay=0.01)
model.compile(loss="sparse_categorical_crossentropy", optimizer=optimizer,
              metrics=["accuracy"])
history = model.fit(train_set, epochs=5, validation_data=valid_set)

```

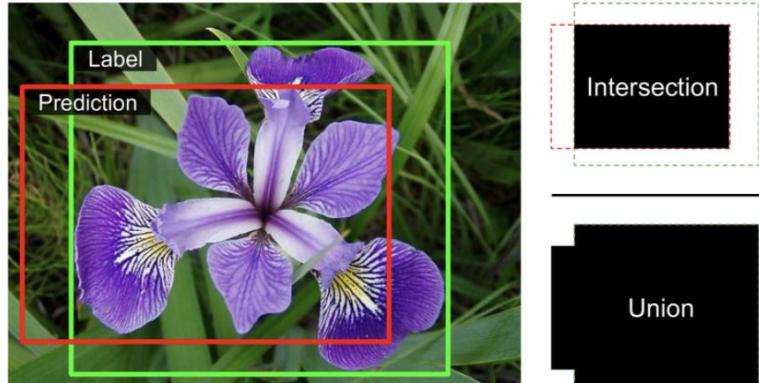


Figure 14-23. Intersection over Union (IoU) metric for bounding boxes

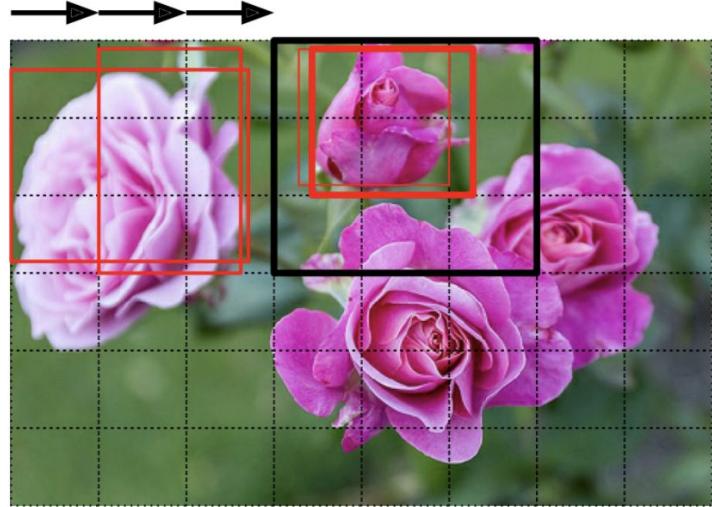


Figure 14-24. Detecting multiple objects by sliding a CNN across the image

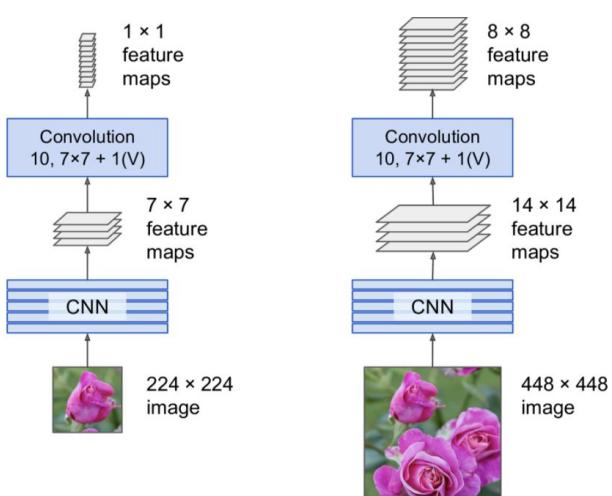
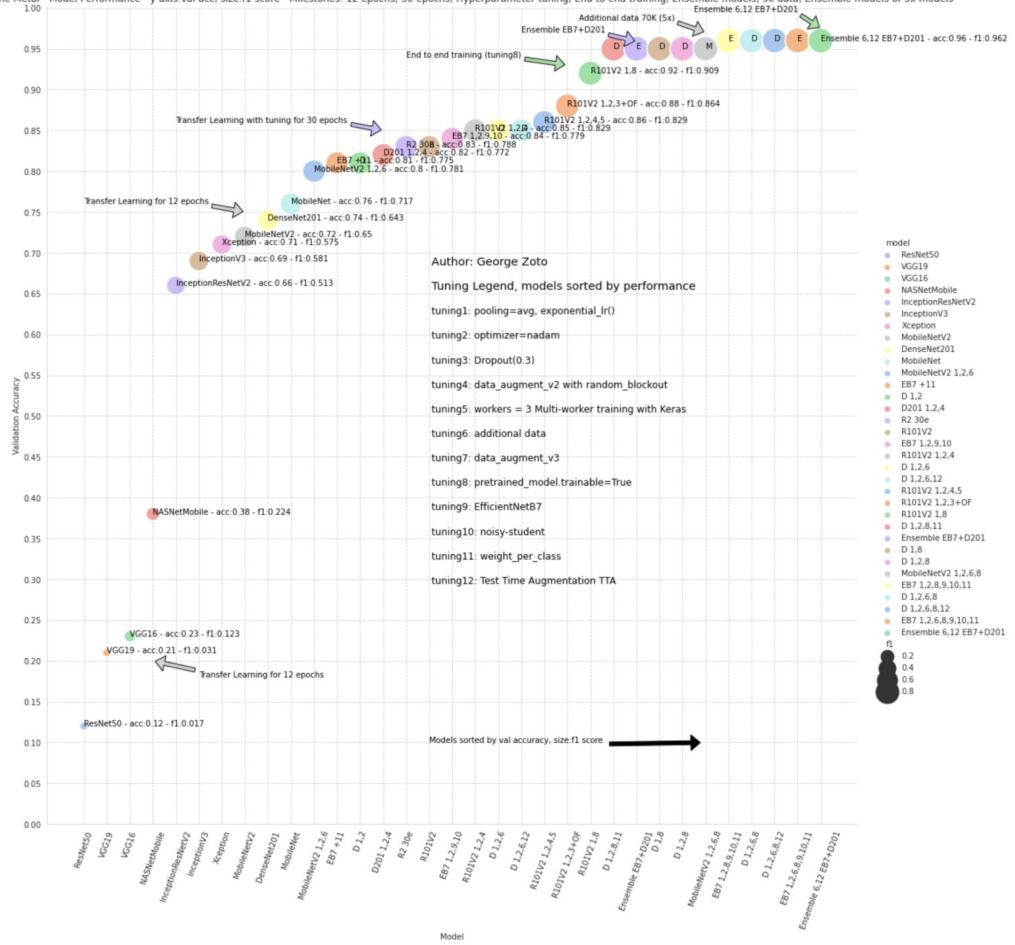


Figure 14-25. The same fully convolutional network processing a small image (left) and a large one (right)



Figure 14-26. Semantic segmentation

Petals to the Metal - Model Performance - y axis:val acc, size:f1 score - Milestones: 12 epochs; 30 epochs; Hyperparameter tuning; End to end training; Ensemble models; 5x data; Ensemble models of 5x models

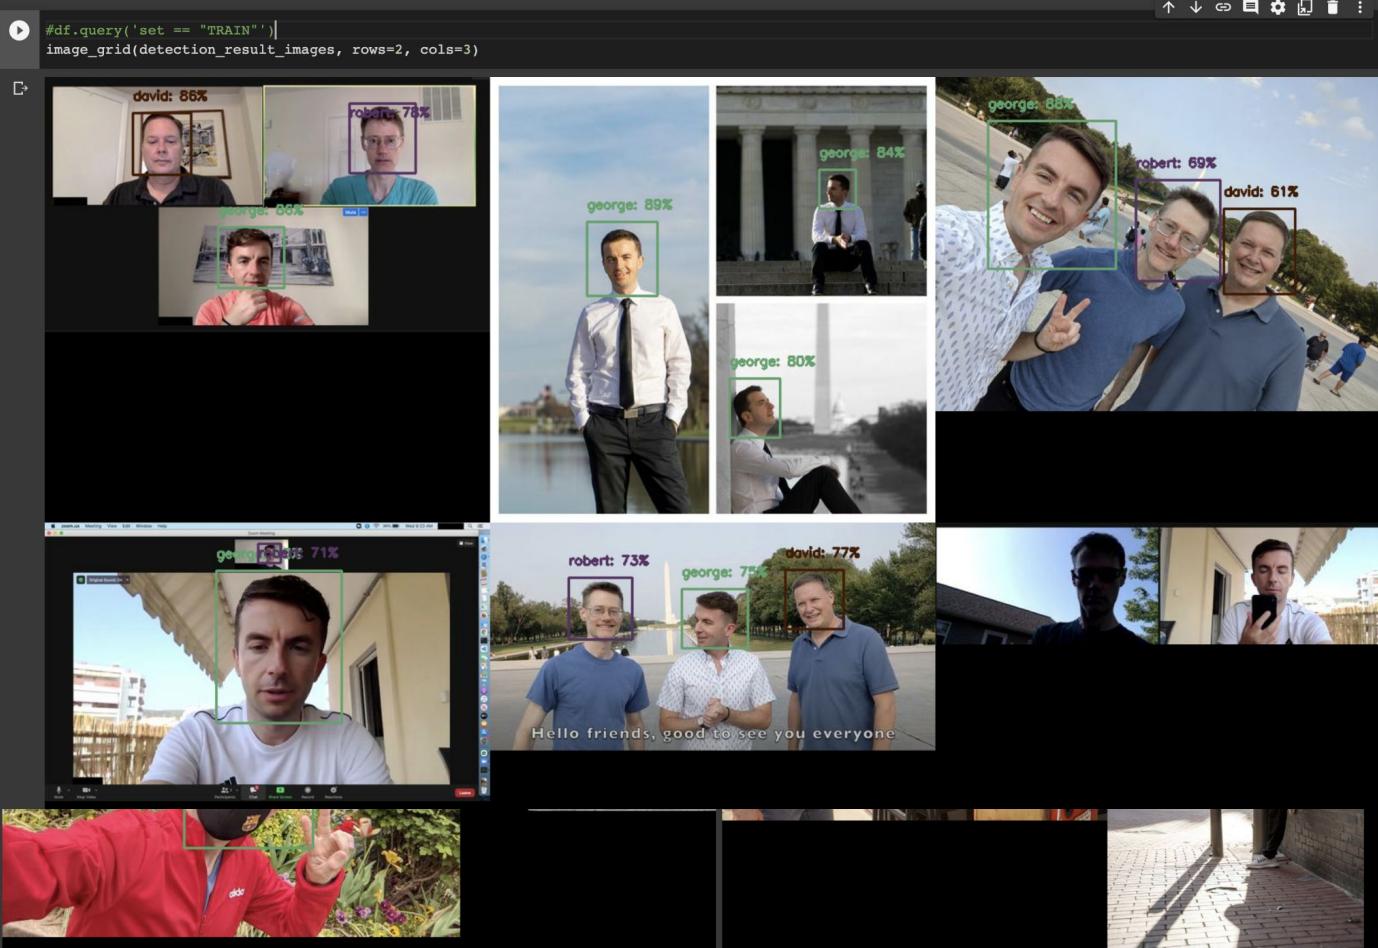


Computer Vision - Petals to the Metal

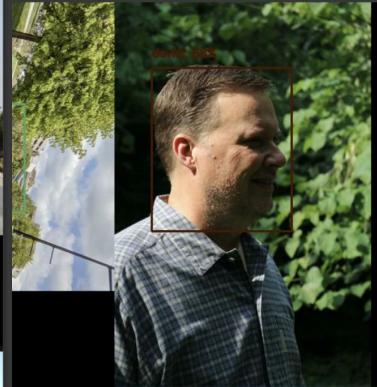


Colab Demo

Deep Learning Adventures



Animated - Best viewed
as a presentation



Deep Learning Adventures

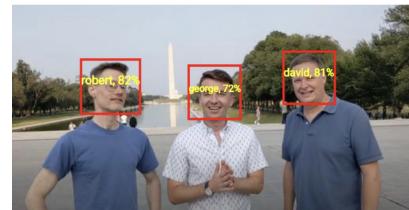
Android App Demo

Deep Learning Adventures

Deep Learning Adventures

Custom Object Detection App

Deep Learning Adventures



Select a preset image or take a new photo



Take photo

Deep Learning Adventures

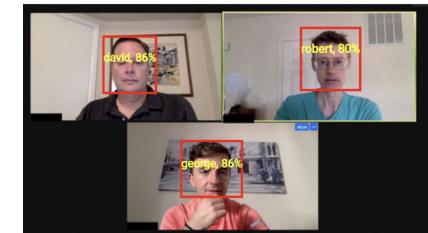


Select a preset image or take a new photo



Take photo

Deep Learning Adventures



Select a preset image or take a new photo



Take photo

Useful Resources

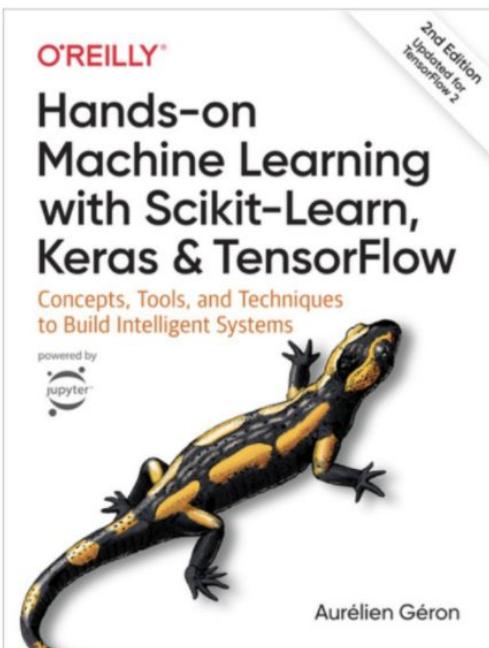
- Long list coming up, see live presentation for now

Deep Learning Adventures + San Diego Machine Learning

Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow, 2nd Edition

★★★★★ [92 reviews](#)

By [Aurélien Géron](#)



TIME TO COMPLETE:

24h 18m

TOPICS:

[Machine Learning](#)

PUBLISHED BY:

[O'Reilly Media, Inc.](#)

PUBLICATION DATE:

September 2019

PRINT LENGTH:

848 pages

Chapter Title

1	The Machine Learning Landscape
2	End-to-End Machine Learning Project
3	Classification
4	Training Models
5	Support Vector Machines
6	Decision Trees
7	Ensemble Learning and Random Forests
8	Dimensionality Reduction
9a	Unsupervised Learning Techniques: Clustering
9b	Unsupervised Learning Techniques: Gaussian Mixtures
10	Introduction to Artificial Neural Networks with Keras
11	Training Deep Neural Networks
12	Custom Models and Training with TensorFlow
13	Loading and Preprocessing Data with TensorFlow
14	Deep Computer Vision Using Convolutional Neural Networks
15	Processing Sequences Using RNNs and CNNs
16	Natural Language Processing with RNNs and Attention
17	Representation Learning and Generative Learning Using Autoencoders and GANs
18	Reinforcement Learning
19	Training and Deploying TensorFlow Models at Scale

Chapter 14 : Deep Computer
Vision Using Convolutional
Neural Networks
Discussion led by George Zoto